

Monitoring





WORKFORCE DEVELOPMENT



Advanced Dashboarding Overview

- You will build a production-grade Grafana dashboard focused on service health.
- The dashboard must support on-call decision making using real metrics and logs.
 - Grafana dashboards
 - Prometheus metrics
 - Loki logs
 - SLO-driven views

Service Scenario

- Your team owns an internal API used by multiple systems.
- When issues occur, dashboards must provide fast answers.
 - Shared dependency
 - High impact
 - On-call usage
 - Operational focus

Service Level Objectives

- SLOs define what success means for the service.
- Dashboards must make SLO violations obvious.
 - Availability $\geq 99\%$
 - p95 latency $\leq 500\text{ms}$
 - Measured continuously
 - Visible in Grafana

Operational Questions

- Every panel must answer a real operational question.
- Charts without purpose add noise.
 - Is it up?
 - Is it fast?
 - Is it scaling?
 - Can I debug it?

Community Dashboards

- Grafana provides many community-built dashboards.
- These are good sources of patterns and ideas.
 - grafana.com
 - Import by ID
 - Editable panels
 - Reusable designs

Importing Dashboards

- Imported dashboards are starting points, not final solutions.
- You can copy panels into your own dashboard.
 - Paste dashboard ID
 - Select data source
 - Explore queries
 - Reuse panels

Creating a New Dashboard

- You will build a custom dashboard tailored to your service.
- This ensures metrics match your SLOs.
 - New dashboard
 - Custom panels
 - Service-specific queries
 - Clear layout

Logs Panel

- Logs provide context when metrics show anomalies.
- Loki allows real-time log searching.
 - Logs panel
 - Search variable
 - Filter by app
 - Time correlation

Replica Count Panel

- Replica count shows scaling and stability.
- Unexpected drops indicate failures.
 - Stat or Gauge
 - Available replicas
 - Expected vs actual
 - Autoscaling insight

Latency Panel

- Latency reflects user experience.
- Percentiles capture worst-case behavior.
 - Time series
 - p95 latency
 - histogram_quantile
 - Trend analysis

Availability Panel

- Availability measures successful request rates.
- It directly reflects the SLO.
 - Success ratio
 - Exclude 5xx
 - Stat panel
 - Thresholds

Error Rate Panel

- Errors highlight failures users experience.
- Tracking errors separately improves detection.
 - 5xx responses
 - Rate over time
 - Alert candidate
 - Incident signal

Throughput Panel

- Throughput shows system load.
- It helps explain latency and errors.
 - Requests per second
 - Traffic spikes
 - Correlation
 - Capacity insight

Dashboard Variables

- Variables make dashboards reusable.
- They reduce duplication across environments.
 - Namespace variable
 - Deployment variable
 - Log filters
 - Dynamic queries

Metrics and Logs Together

- Metrics show what is happening.
- Logs explain why it is happening.
 - Start with metrics
 - Jump to logs
 - Time alignment
 - Root cause

Dashboard Polish

- Clear labeling improves usability.
- Dashboards are shared artifacts.
 - Titles
 - Descriptions
 - Panel grouping
 - Consistent naming

Exporting Dashboards

- Dashboards can be versioned as JSON.
- This supports GitOps workflows.
 - Export JSON
 - Store in Git
 - Re-import
 - Share easily

Sharing Dashboards

- Dashboards can be shared internally or publicly.
- Standardization improves operations.
 - Share links
 - Export files
 - [grafana.com](#)
 - Team reuse

Cluster Dashboard

- Application dashboards are not enough.
- Cluster dashboards show platform health.
 - Node metrics
 - Pod counts
 - Infra errors
 - Capacity overview

On-Call Workflow

- Dashboards guide incident response.
- They should support fast decisions.
 - Check availability
 - Inspect latency
 - Review errors
 - Debug via logs

MCQ 1 – Purpose of Dashboard

- What is the main goal of the dashboard?
 - A. Store logs
 - B. Replace alerts
 - C. Evaluate service health
 - D. Monitor only infra



MCQ 1 – Purpose of Dashboard – Answer

- What is the main goal of the dashboard?
 - A. Store logs
 - B. Replace alerts
 - **C. Evaluate service health**
 - D. Monitor only infra



MCQ 2 – SLO

- What does an SLO represent?
 - A. Alert rule
 - B. Business target
 - C. Log format
 - D. Metric type



MCQ 2 – SLO – Answer

- What does an SLO represent?
 - A. Alert rule
 - **B. Business target**
 - C. Log format
 - D. Metric type



MCQ 3 – p95

- What does p95 latency show?
 - A. Average
 - B. Fastest
 - C. Worst 5%
 - D. Total



MCQ 3 – p95 – Answer

- What does p95 latency show?
 - A. Average
 - B. Fastest
 - **C. Worst 5%**
 - D. Total



MCQ 4 – Logs

- Why use logs?
 - A. Replace metrics
 - B. Explain failures
 - C. Store dashboards
 - D. Generate alerts



MCQ 4 – Logs – Answer

- Why use logs?
 - A. Replace metrics
 - **B. Explain failures**
 - C. Store dashboards
 - D. Generate alerts



MCQ 5 – Variables

- Why use variables?
 - A. Faster queries
 - B. Reduce cost
 - C. Reusable dashboards
 - D. Replace PromQL



MCQ 5 – Variables – Answer

- Why use variables?
 - A. Faster queries
 - B. Reduce cost
 - **C. Reusable dashboards**
 - D. Replace PromQL



Alerting Overview

- Alerts notify engineers when action is required.
- They complement dashboards.
 - Grafana alerting
 - Prometheus metrics
 - Slack notifications
 - On-call workflows

Why Alerting Matters

- Dashboards require manual checking.
- Alerts automate detection.
 - Faster response
 - Reduced downtime
 - Clear signals
 - Operational readiness

Forking the Repo

- You will work in your own fork.
- Git tracks all changes.
 - Fork repo
 - Clone locally
 - Safe experimentation
 - GitOps flow

Slack Webhooks

- Slack receives alert notifications.
- Webhooks enable integration.
 - Slack app
 - Incoming webhook
 - Channel selection
 - Secure URL

MySQL Availability Metric

- mysql_up indicates database health.
- It drives the alert.
 - 1 = up
 - 0 = down
 - Prometheus metric
 - Binary signal

Creating an Alert Rule

- Alerts are defined on panels.
- Rules define failure conditions.
 - mysql_up == 0
 - Evaluate every minute
 - For 1 minute
 - Treat no data as alert

Contact Points

- Contact points define where alerts go.
- Slack is used here.
 - Add contact point
 - Slack type
 - Webhook URL
 - Attach to rule

Simulating Failure

- GitOps changes simulate outages.
- ArgoCD enforces desired state.
 - Remove deployment
 - Commit change
 - ArgoCD prunes
 - Failure occurs

Alert Firing

- Grafana detects the failure.
- Slack receives the alert.
 - Slack message
 - Rule validation
 - End-to-end test
 - Incident flow

Restoring Service

- Restoring Git restores the service.
- Alerts automatically resolve.
 - Revert change
 - ArgoCD redeploys
 - mysql_up returns
 - Alert clears

MCQ 1 – Alerting

- What is alerting for?
 - A. Visualize logs
 - B. Notify engineers
 - C. Store metrics
 - D. Replace dashboards



MCQ 1 – Alerting – Answer

- What is alerting for?
 - A. Visualize logs
 - **B. Notify engineers**
 - C. Store metrics
 - D. Replace dashboards



MCQ 2 – Webhook

- Slack webhook does what?
 - A. Scraps metrics
 - B. Sends alerts
 - C. Deploys apps
 - D. Stores logs



MCQ 2 – Webhook – Answer

- Slack webhook does what?
 - A. Scraps metrics
 - **B. Sends alerts**
 - C. Deploys apps
 - D. Stores logs



MCQ 3 – mysql_up

- mysql_up value 0 means?
 - A. Slow
 - B. Down
 - C. Healthy
 - D. Restart



MCQ 3 – mysql_up – Answer

- mysql_up value 0 means?
 - A. Slow
 - **B. Down**
 - C. Healthy
 - D. Restart



MCQ 4 – ArgoCD

- Removing manifest does what?
 - A. Nothing
 - B. Deletes pod
 - C. Prunes resource
 - D. Stops Prometheus



MCQ 4 – ArgoCD – Answer

- Removing manifest does what?
 - A. Nothing
 - B. Deletes pod
 - **C. Prunes resource**
 - D. Stops Prometheus



MCQ 5 – Resolve

- When does alert resolve?
 - A. Slack restart
 - B. Grafana restart
 - C. MySQL restored
 - D. Prometheus restart



MCQ 5 – Resolve – Answer

- When does alert resolve?
 - A. Slack restart
 - B. Grafana restart
 - **C. MySQL restored**
 - D. Prometheus restart



Exporters Overview

- Exporters expose metrics for apps without native support.
- They bridge legacy systems and Prometheus.
 - Standalone processes
 - Expose /metrics
 - Scrapped by Prometheus
 - Operational overhead

Why Exporters Exist

- Many systems lack Prometheus metrics.
- Exporters fill this gap.
 - MySQL
 - Redis
 - Nginx
 - Legacy systems

Exporter Tradeoffs

- Exporters add complexity.
- They must be maintained.
 - Extra deployment
 - Versioning
 - Security
 - Failure points

Native Metrics Preferred

- Native metrics are simpler and safer.
- Use exporters only when required.
 - Fewer components
 - Better context
 - Lower latency
 - Simpler ops

MySQL Exporter

- The MySQL exporter exposes database internals.
- It runs as a separate service.
 - mysqld-exporter
 - Queries internals
 - Prometheus format
 - Metrics endpoint

Exporter Docs

- Documentation explains configuration.
- Metrics and flags are listed.
 - DockerHub
 - GitHub
 - Env vars
 - Supported metrics

Testing Exporters

- Always verify exporters work.
- Curl the metrics endpoint.
 - Find pod
 - Use netshoot
 - Curl /metrics
 - Inspect output

Netshoot Pod

- Netshoot is a debug container.
- It includes network tools.
 - curl
 - dig
 - tcpdump
 - Cluster debugging

ServiceMonitor

- ServiceMonitors define scraping.
- They are Kubernetes CRDs.
 - Label selectors
 - Scrape interval
 - Declarative config
 - GitOps-friendly

Exporter Workflow

- Exporters fit into the scrape pipeline.
- They behave like any metrics endpoint.
 - App stats
 - Exporter
 - ServiceMonitor
 - Prometheus

MCQ 1 – Exporters

- Purpose of exporters?
 - A. Replace Prometheus
 - B. Expose metrics
 - C. Store logs
 - D. Trigger alerts



MCQ 1 – Exporters – Answer

- Purpose of exporters?
 - A. Replace Prometheus
 - **B. Expose metrics**
 - C. Store logs
 - D. Trigger alerts



MCQ 2 – Endpoint

- Metrics exposed at?
 - A. /health
 - B. /status
 - C. /metrics
 - D. /debug



MCQ 2 – Endpoint – Answer

- Metrics exposed at?
 - A. /health
 - B. /status
 - **C. /metrics**
 - D. /debug



MCQ 3 – MySQL

- MySQL exporter does?
 - A. Backup DB
 - B. Run queries
 - C. Expose metrics
 - D. Manage users



MCQ 3 – MySQL – Answer

- MySQL exporter does?
 - A. Backup DB
 - B. Run queries
 - **C. Expose metrics**
 - D. Manage users



MCQ 4 – ServiceMonitor

- ServiceMonitor role?
 - A. Deploy exporters
 - B. Define scraping
 - C. Visualize data
 - D. Store logs



MCQ 4 – ServiceMonitor – Answer

- ServiceMonitor role?
 - A. Deploy exporters
 - **B. Define scraping**
 - C. Visualize data
 - D. Store logs



MCQ 5 – Best Practice

- When use exporters?
 - A. Always
 - B. Only DBs
 - C. When no native metrics
 - D. Only prod



MCQ 5 – Best Practice – Answer

- When use exporters?
 - A. Always
 - B. Only DBs
 - **C. When no native metrics**
 - D. Only prod



Final Lab Overview

- You will build a centralized Automation Gateway API.
- It exposes business metrics and uses GitOps deployment.
 - Automation API
 - Business metrics
 - Argo CD
 - Grafana proof

Problem Statement

- Automation is often fragmented.
- This lab centralizes control.
 - Ad-hoc scripts
 - No visibility
 - Inconsistent tooling
 - Operational risk

Given Environment

- The cluster and observability stack already exist.
- You build on top of them.
 - EKS
 - LGTM stack
 - Reference API
 - Argo CD ready

Your Objective

- Create automation-gateway.
- Deploy and observe it.
 - FastAPI service
 - Argo CD app
 - Custom metrics
 - REST endpoints

API Endpoints

- The API must expose core endpoints.
- Jobs must run asynchronously.
 - /healthz
 - /v1/jobs
 - /v1/jobs/{id}
 - /metrics

Business Metrics

- Metrics reflect automation outcomes.
- They are not infrastructure metrics.
 - Job counter
 - Duration histogram
 - Queue depth
 - Action labels

Argo CD Deployment

- Use apps-of-apps pattern.
- All manifests live in Git.
 - Deployment
 - Service
 - ServiceMonitor
 - Application CR

Validation Steps

- Verify deployment and metrics.
- End-to-end proof is required.
 - kubectl checks
 - Test API
 - Prometheus scrape
 - Grafana panels

Grafana Dashboard

- Create a minimal dashboard.
- It proves business visibility.
 - Jobs per minute
 - p95 duration
 - Queue depth
 - Live data

Why This Matters

- This is a platform control plane.
- Tools can change without breaking APIs.
 - Stable interfaces
 - Replaceable tools
 - Observable automation
 - Production-ready