

Monitoring





WORKFORCE DEVELOPMENT



Observability in Modern Systems

- Observability is the ability to understand what is happening inside a system by looking at the data it produces. In modern distributed systems, this data must explain behavior without requiring manual debugging or guesswork.
 - Observability is about system introspection
 - Explains behavior without direct access
 - Critical for distributed systems
 - Relies on telemetry from the system

Monitoring vs Observability

- Monitoring tells you whether a system is healthy based on predefined checks and thresholds. Observability tells you what the system is doing internally without needing to predict every failure mode.
 - Monitoring is reactive
 - Observability is exploratory
 - Monitoring answers known questions
 - Observability enables unknown questions

Why Logs and Metrics Are Not Enough

- Logs and metrics describe system behavior, but they lack causality and context. They cannot explain how a single request flows across multiple services.
 - Metrics show trends
 - Logs show events
 - Neither shows request paths
 - Neither explains dependencies

The Distributed Systems Problem

- Modern applications are built from many services that communicate over the network. Failures and latency now emerge from interactions between systems, not single components.
 - Requests span services
 - Failures propagate
 - Latency compounds
 - Root cause becomes unclear

What Is OpenTelemetry

- OpenTelemetry is an open standard for generating and exporting telemetry data from applications. It provides a unified way to collect traces, metrics, and logs across systems.
 - Vendor-neutral standard
 - Open-source CNCF project
 - Cross-language support
 - Exports to any backend

The Three Signals

- OpenTelemetry defines three core observability signals that describe system behavior. Each signal answers a different class of operational questions.
 - Traces show request flows
 - Metrics show numerical behavior
 - Logs show events
 - Signals complement each other

What Is a Trace

- A trace represents the full lifecycle of a single request as it moves through a system. It shows how work is broken down across services and components.
 - One trace per request
 - Spans form a tree
 - Parent-child relationships
 - Latency visible end-to-end

What Is a Span

- A span represents a single operation within a trace. It contains timing information and metadata about the work performed.
 - Start and end time
 - Belongs to one trace
 - Can have children
 - Records errors

Why Tracing Exists

- Tracing exists to explain behavior that cannot be inferred from metrics or logs alone. It provides causality across service boundaries.
 - Explains slow requests
 - Explains failures
 - Explains dependencies
 - Enables root cause analysis

OTEL Is Vendor Neutral

- OpenTelemetry focuses on generating telemetry, not storing or visualizing it. This allows the same instrumentation to work with any observability backend.
 - No vendor lock-in
 - Works with many APM platforms
 - Same code across environments
 - Backends can change freely

Why Standards Matter

- Without standards, every vendor requires custom instrumentation and agents. Standards allow tools to interoperate across organizations and platforms.
 - Avoids proprietary agents
 - Enables tool portability
 - Reduces engineering overhead
 - Improves ecosystem growth

OTEL in Cloud Native Systems

- Cloud-native systems are dynamic, ephemeral, and distributed by design. Observability must adapt automatically as services scale and move.
 - Services scale horizontally
 - Instances are short-lived
 - Infrastructure is abstracted
 - Tracing must be automatic

Telemetry as a Control Plane

- Telemetry acts as a feedback system between applications and operators. It enables real-time decision making based on system behavior.
 - Supports incident response
 - Enables capacity planning
 - Drives automation
 - Feeds reliability engineering

OTEL vs APM Agents

- Traditional Application Performance Monitoring (APM) agents are tightly coupled to specific vendors. OTEL decouples instrumentation from analysis.
 - APM agents are proprietary
 - OTEL is open standard
 - OTEL supports multiple backends
 - Instrumentation is portable

What OTEL Does Not Do

- OpenTelemetry does not store, visualize, or alert on data. It only produces and exports telemetry.
 - No built-in dashboards
 - No built-in alerting
 - No data storage
 - Focuses only on collection

Why Tracing Became Critical

- As systems moved from monoliths to microservices, request paths became invisible. Tracing restores visibility across service boundaries.
 - Monoliths had single stack traces
 - Microservices break call stacks
 - Network hides execution flow
 - Tracing reconstructs execution paths

Tracing vs Logging

- Tracing shows structured request paths, while logging shows unstructured events. They answer different operational questions.
 - Tracing is request-centric
 - Logging is event-centric
 - Tracing shows causality
 - Logging shows symptoms

Tracing vs Metrics

- Metrics aggregate system behavior, while tracing shows individual experiences. Both are required for full observability.
 - Metrics show trends
 - Tracing shows individual journeys
 - Metrics detect anomalies
 - Tracing explains anomalies

OTEL Core Components

- OpenTelemetry is built from several core components that work together to generate and move telemetry. Each component plays a specific role in the observability pipeline.
 - Instrumentation libraries
 - SDKs
 - Collectors
 - Exporters
 - Backends

Instrumentation Libraries

- Instrumentation libraries hook into application frameworks and libraries to generate telemetry. They create spans, metrics, and logs automatically.
 - Framework specific
 - Language specific
 - Often auto-injected
 - Minimal code changes

OTEL SDK

- The SDK is responsible for creating, processing, and exporting telemetry data. It lives inside the application process.
 - Creates tracers and meters
 - Manages span lifecycle
 - Handles sampling
 - Sends data to exporters

OTEL Collector

- The OTEL Collector is a standalone service that receives, processes, and forwards telemetry. It acts as a telemetry gateway.
 - Runs out of process
 - Receives OTLP
 - Can transform data
 - Forwards to backends

Why Use a Collector

- Collectors centralize telemetry processing and reduce coupling between applications and backends. They enable routing and enrichment of data.
 - Decouples apps from vendors
 - Enables sampling control
 - Supports multiple backends
 - Simplifies network topology

OTEL Exporters

- Exporters define how telemetry leaves the system. They translate internal data into backend-specific protocols.
 - OTLP exporter
 - Jaeger exporter
 - Zipkin exporter
 - Vendor specific exporters

OTLP Protocol

- OTLP is the standard protocol used by OpenTelemetry to transmit telemetry. It is designed to be efficient and vendor neutral.
 - HTTP or gRPC
 - Binary encoding
 - Standard across languages
 - Default OTEL protocol

Sampling

- Sampling controls how many traces are collected and stored.
It balances visibility with performance and cost.
 - Always on sampling
 - Probabilistic sampling
 - Tail-based sampling
 - Parent-based sampling

Head vs Tail Sampling

- Head sampling happens when a trace starts, while tail sampling happens after the trace completes. Tail sampling allows smarter decisions.
 - Head sampling is fast
 - Tail sampling sees full trace
 - Tail sampling is more expensive
 - Tail sampling enables filtering

Resource Attributes

- Resource attributes describe the entity producing telemetry. They provide identity and context across systems.
 - Service name
 - Environment
 - Region
 - Version

Context Propagation

- Context propagation allows trace data to flow across service boundaries. It ensures that distributed systems share the same trace.
 - Trace context travels with requests
 - Implemented via headers
 - Works across protocols
 - Required for distributed tracing

Trace Context Headers

- Trace context is usually transmitted through standardized HTTP headers. These headers allow downstream services to join the same trace.
 - traceparent header
 - tracestate header
 - W3C standard
 - Supported across vendors

What Is Baggage

- Baggage is key-value metadata that travels with a trace. It provides business and operational context.
 - Propagates across services
 - Used for custom metadata
 - Not used for metrics
 - Should be kept small

Correlation

- Correlation links traces, metrics, and logs together. It enables operators to move between signals seamlessly.
 - Trace ID in logs
 - Span ID in logs
 - Metrics tagged with trace data
 - Enables cross-signal navigation

Service Graphs

- Service graphs visualize communication between services. They are built automatically from trace data.
 - Shows dependencies
 - Shows request volume
 - Shows error rates
 - Shows latency edges

Trace to Logs

- Trace to logs allows operators to jump from a trace directly into relevant logs. This reduces time to root cause.
 - Uses trace ID
 - Works with Loki
 - Filters logs automatically
 - Improves debugging speed

Trace to Metrics

- Trace to metrics allows operators to correlate a single trace with aggregate system behavior. It connects micro and macro views.
 - Uses span attributes
 - Generates RED metrics
 - Supports SLOs
 - Supports capacity analysis

RED Method

- The RED method is a common way to monitor services using request data. It focuses on user experience.
 - Rate
 - Errors
 - Duration
 - Derived from traces

USE Method

- The USE method focuses on resource utilization. It complements RED by describing infrastructure health.
 - Utilization
 - Saturation (CPU queue, swap space)
 - Errors
 - Infrastructure focused

OTEL in Incident Response

- OTEL provides the fastest path from symptom to root cause during incidents. It replaces guesswork with evidence.
 - Find failing service
 - Follow request path
 - Inspect errors
 - Confirm fix with live traces

OTEL and the LGTM Stack

- OpenTelemetry acts as the data plane for the LGTM stack. It feeds traces into Tempo, logs into Loki, and metrics into Prometheus or Mimir.
 - OTEL generates telemetry
 - Tempo stores traces
 - Loki stores logs
 - Prometheus or Mimir stores metrics

Tempo in the LGTM Stack

- Tempo is a distributed tracing backend designed for scale. It stores traces and integrates directly with Grafana.
 - Receives OTLP
 - Stores trace data
 - No indexing by default
 - Optimized for cost

Why Tempo Does Not Index

- Tempo avoids heavy indexing to reduce cost and complexity.
It relies on Grafana for querying and navigation.
 - Lower storage cost
 - Faster ingestion
 - Relies on metadata
 - Optimized for large volumes

Grafana as the Control Plane

- Grafana acts as the single interface for all observability data.
It connects traces, logs, and metrics.
 - Unified dashboards
 - Explore view
 - Trace navigation
 - Signal correlation

Trace Exploration in Grafana

- Grafana allows interactive exploration of traces. Engineers can drill into spans and follow dependencies.
 - Search by service
 - Search by duration
 - View span trees
 - Inspect errors

Trace to Logs in LGTM

- LGTM enables jumping from a trace directly into Loki logs. This connects request paths with raw events.
 - Trace ID used as filter
 - Automatic log queries
 - Reduces investigation time
 - Improves debugging flow

Trace to Metrics in LGTM

- LGTM connects trace attributes with metrics. This enables SLOs and service level monitoring.
 - RED metrics from spans
 - Latency percentiles
 - Error rates
 - Throughput

Service Graphs in LGTM

- Service graphs visualize dependencies using trace data. They help understand system topology.
 - Auto-generated
 - Shows call relationships
 - Shows traffic volume
 - Highlights bottlenecks

OTEL Data Flow

- OTEL defines how telemetry flows from applications to backends. This pipeline must be reliable and observable.
 - App generates telemetry
 - SDK processes data
 - Collector receives data
 - Backends store data

Why Platform Teams Care

- Platform teams use OTEL to provide observability as a service. It becomes shared infrastructure.
 - Standardized telemetry
 - Reduced onboarding time
 - Consistent tooling
 - Improved reliability

OTEL in FastAPI

- FastAPI can be instrumented with OpenTelemetry to automatically generate traces for every request. This requires no changes to application logic.
 - Auto-instrumentation supported
 - Works with ASGI middleware
 - Creates spans per endpoint
 - Captures request metadata

Auto-Instrumentation

- Auto-instrumentation injects telemetry into libraries at runtime. It eliminates the need to manually create spans for basic operations.
 - Uses monkey patching
 - Configured at startup
 - Minimal code changes
 - Covers common frameworks

opentelemetry-instrument

- The opentelemetry-instrument command wraps the application process. It bootstraps the OTEL SDK before the app starts.
 - CLI wrapper
 - Initializes providers
 - Loads instrumentation packages
 - Requires correct dependencies

OTLP Endpoints

- OTLP endpoints define where telemetry is sent. They must be reachable from the application environment.
 - HTTP or gRPC
 - Usually port 4317 or 4318
 - Collector or backend
 - Critical for connectivity

Custom Spans

- Custom spans allow developers to trace business logic. They provide visibility into internal operations.
 - Manually defined
 - Used for critical code paths
 - Add custom attributes
 - Expose performance bottlenecks

Error Tracing

- Unhandled exceptions automatically mark spans as errors.
This allows traces to represent failures.
 - Exception recorded in span
 - Status set to error
 - Searchable in Tempo
 - Supports incident analysis

Outbound Calls

- Outbound HTTP calls can also be traced if instrumentation exists. This enables full request graphs.
 - HTTP client instrumentation
 - Database instrumentation
 - Message queue instrumentation
 - Shows dependency latency

Sampling in Production

- Sampling strategies must balance cost and visibility.
Production systems rarely sample 100 percent.
 - Lower cost
 - Reduced storage
 - Focus on errors
 - Tail-based filtering

Performance Overhead

- OTEL introduces minimal overhead when configured correctly. Overhead increases with sampling and export volume.
 - CPU overhead
 - Network overhead
 - Storage overhead
 - Generally acceptable in prod

Common Production Patterns

- OTEL is typically deployed with a collector and centralized configuration. This simplifies management and scaling.
 - Sidecar collectors
 - DaemonSet collectors
 - Central sampling rules
 - Multi-backend routing

MCQ 1

- What is the main purpose of OpenTelemetry?
- A. Store logs and metrics
- B. Provide dashboards and alerts
- C. Generate and export telemetry
- D. Replace Prometheus

MCQ 1 – Answer

- What is the main purpose of OpenTelemetry?
- A. Store logs and metrics
- B. Provide dashboards and alerts
- **C. Generate and export telemetry**
- D. Replace Prometheus

MCQ 2

- Which component is responsible for receiving and forwarding telemetry out of process?
- A. SDK
- B. Instrumentation library
- C. Collector
- D. Exporter

MCQ 2 – Answer

- Which component is responsible for receiving and forwarding telemetry out of process?
- A. SDK
- B. Instrumentation library
- ****C. Collector****
- D. Exporter

MCQ 3

- What does context propagation enable?
- A. Log aggregation
- B. Metric sampling
- C. Distributed tracing across services
- D. Alert routing

MCQ 3 – Answer

- What does context propagation enable?
- A. Log aggregation
- B. Metric sampling
- **C. Distributed tracing across services**
- D. Alert routing

MCQ 4

- Why is tail-based sampling powerful?
- A. It samples faster
- B. It reduces CPU usage
- C. It sees the full trace before deciding
- D. It eliminates exporters

MCQ 4 – Answer

- Why is tail-based sampling powerful?
- A. It samples faster
- B. It reduces CPU usage
- **C. It sees the full trace before deciding**
- D. It eliminates exporters

MCQ 5

- Which LGTM component stores distributed traces?
- A. Loki
- B. Prometheus
- C. Grafana
- D. Tempo

MCQ 5 – Answer

- Which LGTM component stores distributed traces?
- A. Loki
- B. Prometheus
- C. Grafana
- **D. Tempo**