

Programming for Automation





WORKFORCE DEVELOPMENT



Why APIs Need Authentication

- Most real-world APIs are not publicly accessible and require authentication. Authentication allows servers to identify who is making a request.
 - Prevents unauthorized access
 - Enables per-user permissions
 - Allows usage tracking and auditing

What Is an API Key?

- An API key is a long, random value used to authenticate API requests. It is sent by the client with each request to prove identity.
 - Acts like a password for programs
 - Usually generated by the server
 - Must be kept secret

How API Keys Are Sent

- API keys are typically sent in HTTP headers rather than URLs. Headers avoid leaking secrets into logs and browser history.
 - Authorization: Bearer <token>
 - X-API-Key: <token>
 - Avoid query parameters for secrets

API Keys in the Real World

- Most cloud and SaaS platforms rely on API keys for automation access. Keys are designed for scripts, CI/CD, and system-to-system communication.
 - Used by AWS, GitHub, Docker, and others
 - Common in CLI tools and pipelines
 - Often permission-scoped

Security Expectations

- API keys must be treated with the same care as passwords. Poor key handling leads directly to security breaches.
 - Never commit keys to source control
 - Rotate keys regularly
 - Revoke compromised keys immediately

Custom Headers vs Authorization

- APIs commonly accept keys using the Authorization header or a custom header. Both approaches work, but Authorization is more standardized.
 - Authorization: Bearer <key>
 - X-API-Key: <key>
 - Be consistent across endpoints

Why Not Put Keys in URLs

- URLs are logged, cached, and stored in browser history. Placing secrets in URLs increases the risk of accidental exposure.
 - URLs appear in logs
 - Browsers store URL history
 - Headers are safer for secrets

Treat API Keys Like Passwords

- API keys grant access to protected resources. Anyone with the key can act as the associated user or service.
 - Never share keys publicly
 - Rotate keys periodically
 - Revoke keys when compromised

Randomness Matters

- API keys must be difficult to guess. Predictable keys are vulnerable to brute-force attacks.
 - Use secrets module
 - Avoid sequential or short keys
 - Prefer long random values

Key Rotation

- Key rotation limits the damage of leaked credentials. Good systems support replacing keys without downtime.
 - Generate new key
 - Update clients
 - Revoke old key

Key Revocation

- Revocation allows immediate removal of access. This is critical when a key is exposed or abused.
 - Disable or delete key
 - Return 401 for revoked keys
 - Log revocation events

What API Keys Do

- API keys identify the caller and allow access decisions. They enable rate limits, permissions, and auditing.
 - Identify who is calling
 - Control access scope
 - Support monitoring

What API Keys Do Not Do

- API keys do not encrypt network traffic. They rely on HTTPS for confidentiality.
 - No encryption by themselves
 - Visible on the wire without HTTPS
 - Must be combined with TLS

HTTPS Is Still Required

- Without HTTPS, API keys can be intercepted in transit. Encryption is mandatory for any authenticated API.
 - Protects against sniffing
 - Prevents man-in-the-middle attacks
 - Required for secure cookies

Storing API Keys Safely

- Storing raw API keys creates unnecessary risk. Production systems store only derived representations.
 - Avoid plaintext storage
 - Use hashing techniques
 - Limit blast radius of leaks

Why Hash API Keys

- Storing raw API keys creates a single point of failure. Hashing ensures leaked databases cannot be used directly for access.
 - Protects against database leaks
 - Industry best practice
 - Same principle as password storage

What Is a Hash Function

- A hash function converts input data into a fixed-length output. The process is one-way and deterministic.
 - Same input always produces same output
 - Output length is fixed
 - Original value cannot be recovered

One-Way by Design

- Hash functions are intentionally irreversible. This prevents attackers from recovering secrets from stored hashes.
 - No decryption step exists
 - Verification is done by comparison
 - Security relies on math

SHA-256 Overview

- SHA-256 is a cryptographic hash function standardized and widely trusted. It produces a 256-bit output suitable for security applications.
 - Collision resistant
 - Deterministic
 - Fast enough for API keys

Why Python hash() Is Not Secure

- Python's built-in hash() is designed for hash tables, not security. Its output is intentionally unstable across runs.
 - Changes between executions
 - Not cryptographically secure
 - Never store security data with it

Hashing API Keys in Practice

- When an API key is presented, the server hashes it before comparison. The raw key is never stored or logged.
 - Store only hashes
 - Hash on every request
 - Compare hash values

What Is HMAC

- HMAC combines a hash function with a server-side secret. This prevents offline guessing even if hashes are leaked.
 - Uses a shared secret
 - Secret is not stored in database
 - Stronger than plain hashing

Why Use HMAC for API Keys

- API keys are long and random, making HMAC practical and secure. Attackers cannot validate guesses without the server secret.
 - Protects against brute-force
 - Secret stored separately
 - Industry-recommended approach

HMAC vs Salting

- Salts are stored with hashes, while HMAC secrets are not. This changes the attacker's threat model significantly.
 - HMAC secret is server-only
 - Salts are public
 - HMAC blocks offline verification

Secure Comparison

- Comparing hashes must be done carefully. Naive string comparison can leak timing information.
 - Use `hmac.compare_digest()`
 - Avoid `==` for secrets
 - Prevents timing attacks

Generating API Keys Securely

- API keys must be long, random, and unguessable. Python provides a safe standard library for this purpose.
 - Use `secrets.token_hex()`
 - Avoid predictable values
 - Generate once, show once

Key Length Matters

- Short keys are easier to brute-force. Long random keys make guessing computationally infeasible.
 - Minimum 128 bits recommended
 - Hex encoding is common
 - Random beats complexity

Where API Keys Live

- Clients store API keys outside of source code. Servers never expose keys after creation.
 - Environment variables
 - Secret managers
 - CI/CD secrets

Authorization Header Flow

- Most APIs expect keys in the Authorization header. This keeps credentials out of URLs and logs.
 - Authorization: Bearer <key>
 - Standardized pattern
 - Works with proxies

Custom API Key Headers

- Some APIs use custom headers for clarity or tooling reasons.
Security properties remain the same.
 - X-API-Key header
 - Easier for scripts
 - Still requires HTTPS

Authentication vs Authorization

- Authentication identifies who is calling the API. Authorization determines what they are allowed to do.
 - Key validates identity
 - Role controls access
 - Separate concerns

Correct Error Responses

- Clients rely on HTTP status codes to understand failures.
Authentication errors must be explicit and consistent.
 - 401 for missing or invalid keys
 - 403 for insufficient permissions
 - Avoid leaking details

Why Error Messages Matter

- Overly descriptive errors help attackers. Good errors help developers without exposing secrets.
 - Do not reveal valid usernames
 - Do not echo keys
 - Be precise but vague

Logging Authentication Attempts

- Authentication events are security-relevant signals. Logs help detect abuse and misconfiguration.
 - Log failures
 - Track source IP
 - Avoid logging secrets

Lab 1 Summary

- API keys are simple but powerful when handled correctly.
Security depends on storage, transport, and validation.
 - Headers not URLs
 - Hash and HMAC
 - HTTPS always

Which header is most commonly used to send API keys?

- A. Cookie
- B. Authorization
- C. Content-Type
- D. Host



Why should API keys not be placed in URLs?

- A. URLs cannot carry headers
- B. URLs are slower
- C. URLs are logged and cached
- D. Browsers block them



What is the main purpose of hashing API keys?

- A. Compress keys
- B. Encrypt traffic
- C. Protect against database leaks
- D. Improve performance



Which Python function is safest for comparing API key hashes?

- A. ==
- B. equals()
- C. hmac.compare_digest()
- D. str.compare()



Which status code should be returned for a missing API key?

- A. 400
- B. 401
- C. 403
- D. 404



Answer: Which header is most commonly used to send API keys?

- A. Cookie
- **B. Authorization**
- C. Content-Type
- D. Host



Answer: Why should API keys not be placed in URLs?

- A. URLs cannot carry headers
- B. URLs are slower
- **C. URLs are logged and cached**
- D. Browsers block them



Answer: What is the main purpose of hashing API keys?

- A. Compress keys
- B. Encrypt traffic
- **C. Protect against database leaks**
- D. Improve performance



Answer: Which Python function is safest for comparing API key hashes?

- A. ==
- B. equals()
- **C. hmac.compare_digest()**
- D. str.compare()



Answer: Which status code should be returned for a missing API key?

- A. 400
- **B. 401**
- C. 403
- D. 404



Lab 2: JWT vs Session Authentication

- This lab compares traditional session-based authentication with modern JWT-based approaches. Both solve the same problem but scale very differently.
 - Authentication model comparison
 - Stateful vs stateless
 - Real-world usage patterns

What Is Authentication State

- Authentication state tracks who a user is across requests.
Different systems store this state in different places.
 - Session IDs
 - Cookies
 - Tokens

Session-Based Authentication

- Traditional web apps rely on server-managed session identifiers.
The server must remember every active user session.
 - Random session ID
 - Stored server-side
 - Mapped to user

How Sessions Work

- After login, the server creates a session and sends the ID to the client. Every request must be validated against session storage.
 - Session ID in cookie
 - Lookup on each request
 - Stateful design

Session Storage Options

- Sessions can be stored in memory or external systems. Scaling requires shared session storage.
 - In-memory
 - Redis
 - Databases

Problems With Sessions

- Sessions increase server complexity and coupling. Scaling horizontally becomes harder.
 - Sticky sessions
 - Extra infrastructure
 - Operational overhead

Introducing JWT

- JSON Web Tokens embed authentication claims inside the token itself. The server does not store per-user session state.
 - Self-contained token
 - Signed not encrypted
 - Stateless

What JWT Contains

- JWTs carry claims about the user and expiration time. Anyone can read the payload, but only the server can sign it.
 - Username
 - Expiration
 - Signature

JWT Verification Flow

- Each request presents the token to the server. The server validates signature and expiration.
 - No database lookup
 - Signature verification
 - Time-based expiry

JWT vs Sessions Summary

- JWTs trade revocation simplicity for scalability. Sessions trade scalability for control.
 - Stateless vs stateful
 - Scale vs revoke
 - Choose by use case

JWT Issuance Flow

- JWTs are created after successful authentication. They represent a signed statement from the server.
 - Issued at login
 - Signed with secret
 - Contains claims

JWT Claims

- Claims are pieces of information stored inside the token. They describe identity and validity.
 - username
 - expiration (exp)
 - optional permissions

JWT Expiration

- JWTs always include an expiration timestamp. Expired tokens must be rejected.
 - Limits damage if stolen
 - Encourages re-authentication
 - Time-based control

Why Short-Lived Tokens

- JWTs cannot be instantly revoked. Short lifetimes reduce risk.
 - Minutes not days
 - Rotate often
 - Delete cookie on logout

JWT Transport Options

- JWTs can be sent via headers or cookies. Cookies are preferred for browsers.
 - Authorization header
 - HttpOnly cookie
 - SameSite support

HttpOnly Cookies

- HttpOnly cookies cannot be accessed by JavaScript. This protects tokens from XSS.
 - Browser-managed
 - Not readable via JS
 - Safer default

SameSite Cookies

- SameSite controls cross-site cookie behavior. It reduces CSRF risk.
 - Strict
 - Lax
 - None

Secure Cookie Flag

- Secure cookies are only sent over HTTPS. They prevent token leakage on plaintext connections.
 - Requires HTTPS
 - Protects credentials
 - Mandatory in production

JWT Validation Errors

- JWT verification can fail for multiple reasons. Each failure must be handled safely.
 - Missing token
 - Expired token
 - Invalid signature

JWT Logout Behavior

- Logging out deletes the JWT cookie. The token itself remains valid until expiration.
 - Client-side removal
 - No server revocation
 - Expiration enforced

JWT Is Not Encryption

- JWT payloads are base64 encoded, not encrypted. Anyone can read the contents if they have the token.
 - Readable payload
 - Signature enforces trust
 - Do not store secrets inside

JWT Signature Purpose

- The signature proves the token was created by the server.
Clients cannot forge valid tokens.
 - Server-only secret
 - Tamper detection
 - Trust boundary

What Happens If a JWT Is Stolen

- A stolen JWT is valid until it expires. This is why short lifetimes are critical.
 - No instant revocation
 - Expiration limits damage
 - Logout deletes cookie only

CSRF and JWT Cookies

- Cookies are sent automatically by browsers. CSRF protections are still required.
 - SameSite cookies
 - CSRF tokens
 - Origin checks

JWT in Authorization Headers

- APIs sometimes send JWTs via headers instead of cookies. This is common for mobile and service clients.
 - Authorization: Bearer
 - Manual attachment
 - Different threat model

JWT for Browsers vs APIs

- Browser-based apps and API clients have different needs.
Transport choice matters.
 - Cookies for browsers
 - Headers for services
 - Security tradeoffs

Stateless Scaling Benefits

- JWTs allow horizontal scaling without shared session stores. Any server can validate any token.
 - No Redis needed
 - No sticky sessions
 - Simpler infrastructure

JWT Operational Tradeoffs

- Stateless design moves complexity elsewhere. Expiration and key rotation become critical.
 - Rotate secrets carefully
 - Short expirations
 - Plan revocation strategy

JWT Rotation

- JWT signing secrets must be rotated periodically. Old tokens become invalid when secrets change.
 - Key rotation events
 - Forced logout
 - Operational planning

Lab 2 Summary

- JWTs replace server-side sessions with signed tokens. They scale well but require careful security design.
 - Stateless auth
 - Short-lived tokens
 - Cookie security

What makes JWT authentication stateless?

- A. Cookies store sessions
- B. Tokens store claims
- C. Server stores sessions
- D. Redis is optional



Why are JWTs usually short-lived?

- A. They are encrypted
- B. They are cached
- C. They cannot be revoked
- D. They are compressed



Which cookie flag prevents JavaScript access?

- A. Secure
- B. SameSite
- C. HttpOnly
- D. Path



What happens when a JWT signing secret is rotated?

- A. Tokens are refreshed
- B. Old tokens become invalid
- C. Cookies persist
- D. Sessions migrate



Which risk still exists when using JWT cookies?

- A. SQL injection
- B. XSS
- C. CSRF
- D. Brute force



Answer: What makes JWT authentication stateless?

- A. Cookies store sessions
- **B. Tokens store claims**
- C. Server stores sessions
- D. Redis is optional



Answer: Why are JWTs usually short-lived?

- A. They are encrypted
- B. They are cached
- **C. They cannot be revoked**
- D. They are compressed



Answer: Which cookie flag prevents JavaScript access?

- A. Secure
- B. SameSite
- **C. HttpOnly**
- D. Path



Answer: What happens when a JWT signing secret is rotated?

- A. Tokens are refreshed
- **B. Old tokens become invalid**
- C. Cookies persist
- D. Sessions migrate



Answer: Which risk still exists when using JWT cookies?

- A. SQL injection
- B. XSS
- **C. CSRF**
- D. Brute force



Lab 3: Securing APIs with HTTPS

- HTTPS encrypts data between clients and servers. It is mandatory for protecting credentials and tokens.
 - Encryption
 - Authentication
 - Integrity

What Is HTTPS

- HTTPS is HTTP layered on top of TLS. It prevents attackers from reading or modifying traffic.
 - Encrypted transport
 - Identity verification
 - Tamper resistance

Why HTTP Is Dangerous

- Plain HTTP sends data in cleartext. Anyone on the network can intercept credentials.
 - Passwords exposed
 - Tokens leaked
 - Man-in-the-middle attacks

TLS in Simple Terms

- TLS creates a secure tunnel before any data is exchanged. All application data flows through this encrypted channel.
 - Handshake
 - Key exchange
 - Secure session

Public and Private Keys

- TLS relies on asymmetric cryptography. Public keys are shared while private keys remain secret.
 - Public key advertised
 - Private key protected
 - Trust model

Certificates

- Certificates bind a public key to a server identity. Clients trust certificates signed by known authorities.
 - Contains public key
 - Signed by CA
 - Proves identity

Digital Signatures

- Servers prove ownership of private keys during TLS handshakes.
Clients verify signatures using public keys.
 - Authentication
 - Integrity
 - Non-repudiation

Encryption After Handshake

- Once TLS is established, symmetric encryption is used. This keeps communication fast and secure.
 - Session keys
 - Efficient encryption
 - Secure channel

HTTPS and Cookies

- Secure cookies are only sent over HTTPS. Authentication cookies require HTTPS.
 - secure flag
 - HttpOnly flag
 - SameSite flag

Lab 3 Scope

- This lab focuses on enabling HTTPS locally and understanding TLS. Production certificate management is out of scope.
 - Local development
 - Self-signed certs
 - Conceptual grounding

Self-Signed Certificates

- Local development often uses self-signed certificates. These provide encryption but not public trust.
 - Encrypted traffic
 - Browser warnings expected
 - Dev-only usage

Why Browsers Warn

- Browsers trust certificates signed by known authorities. Self-signed certificates cannot be verified.
 - No trusted CA
 - Manual acceptance required
 - Expected behavior

OpenSSL Overview

- OpenSSL is a command-line tool for cryptographic operations.
It is commonly used to generate TLS certificates.
 - Key generation
 - Certificate creation
 - Industry standard

Generating a Private Key

- TLS starts with generating a private key. This key must be kept secret.
 - openssl genrsa
 - 2048-bit minimum
 - Never share

Creating a Certificate

- Certificates bind a public key to an identity. Self-signed certs are sufficient for learning.
 - Certificate signing request
 - Self-signing
 - Expiration dates

Certificate Files

- Flask requires both a certificate and a private key. These files enable HTTPS locally.
 - cert.pem
 - key.pem
 - Filesystem permissions

Flask HTTPS Configuration

- Flask can serve HTTPS using an SSL context. This enables encrypted local testing.
 - `ssl_context` parameter
 - cert and key paths
 - Debug mode supported

Secure Cookies with HTTPS

- Cookies marked Secure are only sent over HTTPS. Browsers enforce this automatically.
 - secure=True
 - Required for auth cookies
 - HTTPS-only transport

Testing HTTPS with curl

- curl can test HTTPS endpoints locally. Self-signed certs require explicit trust bypass.
 - curl -k
 - Dev-only flag
 - Never use in production

Lab 3 Summary

- HTTPS protects credentials and tokens in transit. Authentication is unsafe without encryption.
 - TLS encryption
 - Secure cookies
 - Production requirement

What does HTTPS primarily protect?

- A. Application logic
- B. Data in transit
- C. Server uptime
- D. Database encryption



Why are self-signed certificates acceptable in labs?

- A. They are faster
- B. They provide encryption
- C. Browsers trust them
- D. They scale better



Which file must remain secret on the server?

- A. cert.pem
- B. csr.pem
- C. key.pem
- D. openssl.conf



Why will browsers refuse Secure cookies over HTTP?

- A. Performance reasons
- B. Missing headers
- C. Encryption required
- D. Cookie size limits



What does curl -k do?

- A. Encrypts traffic
- B. Enables cookies
- C. Ignores cert trust errors
- D. Generates certificates



Answer: What does HTTPS primarily protect?

- A. Application logic
- **B. Data in transit**
- C. Server uptime
- D. Database encryption



Answer: Why are self-signed certificates acceptable in labs?

- A. They are faster
- **B. They provide encryption**
- C. Browsers trust them
- D. They scale better



Answer: Which file must remain secret on the server?

- A. cert.pem
- B. csr.pem
- **C. key.pem**
- D. openssl.conf



Answer: Why will browsers refuse Secure cookies over HTTP?

- A. Performance reasons
- B. Missing headers
- **C. Encryption required**
- D. Cookie size limits



Answer: What does curl -k do?

- A. Encrypts traffic
- B. Enables cookies
- **C. Ignores cert trust errors**
- D. Generates certificates

