

Programming for Automation





WORKFORCE DEVELOPMENT



API Rate Limiting

- Rate limiting controls how frequently clients can call an API over time. It protects services from abuse while ensuring fair access for all users.
 - Limits requests per client or API key
 - Prevents denial-of-service attacks
 - Protects backend resources

Why Rate Limiting Matters

- Without limits, a single client can overwhelm your API and impact others. Rate limiting enforces fairness and stability.
 - Avoids resource exhaustion
 - Controls unexpected traffic spikes
 - Keeps APIs reliable

HTTP 429 Too Many Requests

- When a client exceeds the allowed request rate, the API should reject the request. HTTP 429 signals the client to slow down.
 - Returned when rate limit is exceeded
 - Client request is valid but not allowed yet
 - Used instead of generic 400 errors

Retry-After Header

- The Retry-After header tells the client how long to wait before retrying. This enables polite, predictable client behavior.
 - Sent with HTTP 429 responses
 - Value is time in seconds
 - Improves client reliability

Rate Limiting Scope

- Rate limits are usually applied per client identity, not globally. API keys are commonly used as the identifier.
 - Per API key
 - Per IP address
 - Per user account

Where Rate Limiting Lives

- Rate limiting can exist at multiple layers of a system. This lab focuses on application-level rate limiting.
 - Application code
 - API gateways
 - Reverse proxies

In-Memory Rate Limiting

- This lab uses in-memory data structures to track request rates. This is simple but limited to a single process.
 - Easy to implement
 - Resets on restart
 - Not shared across servers

Common Algorithms

- Different algorithms balance fairness, burst handling, and memory usage. Two common ones are sliding window and token bucket.
 - Sliding window log
 - Token bucket
 - Fixed window (not used here)

Sliding Window Concept

- Sliding window tracks exact request timestamps within a rolling time window. It provides precise rate enforcement.
 - Tracks recent request times
 - Removes expired timestamps
 - Fair but memory-heavy

Token Bucket Concept

- Token bucket allows controlled bursts while enforcing a long-term average rate. Tokens refill over time.
 - Each request consumes a token
 - Tokens refill gradually
 - Allows short bursts

Client-Side Behavior

- Rate limiting requires cooperation from both server and client.
Clients should slow down instead of retrying immediately.
 - Read Retry-After header
 - Wait before retrying
 - Avoid hammering the API

Idempotent Requests

- Retries are safest for requests that do not change server state. GET requests are typically idempotent.
 - GET is safe to retry
 - POST can cause duplicates
 - PUT requires caution

Sliding Window Log

- Sliding window log enforces limits by tracking exact request timestamps. It gives precise control over request rates.
 - Uses timestamps per request
 - Maintains rolling window
 - Accurate but memory-heavy

Sliding Window Tradeoffs

- This approach is fair but can consume memory for busy clients. It is best for low to moderate traffic.
 - Stores many timestamps
 - Per-key tracking
 - Not ideal at massive scale

Token Bucket Model

- Token bucket smooths traffic while allowing short bursts. It is commonly used in production systems.
 - Tokens refill over time
 - Each request consumes a token
 - Allows controlled bursts

Token Bucket Tradeoffs

- Token bucket uses less memory and supports bursts. It is slightly less precise than sliding window.
 - Lower memory usage
 - Burst-friendly
 - Approximate enforcement

Choosing an Algorithm

- The right algorithm depends on traffic patterns and fairness requirements. Both approaches are valid.
 - Sliding window = precision
 - Token bucket = flexibility
 - Gateways often choose token bucket

Where to Enforce Limits

- Rate limiting can be enforced at different layers. This lab demonstrates application-level control.
 - Application code
 - Reverse proxy
 - API gateway

Distributed Systems

- In-memory limits do not work across multiple servers. Production systems require shared state.
 - Redis-backed counters
 - Centralized gateways
 - Consistent enforcement

Lab Focus

- This lab keeps rate limiting simple to focus on concepts. The same ideas apply at larger scale.
 - Per API key limits
 - HTTP 429 responses
 - Retry-After handling

Implementing Sliding Window

- Sliding window uses timestamps to decide whether a request is allowed. Old timestamps are removed before evaluating the limit.
 - Track requests per API key
 - Remove expired timestamps
 - Allow or deny request

Time Sources

- Using a monotonic clock avoids issues with system clock changes. This makes rate limiting more reliable.
 - `time.monotonic()`
 - Unaffected by NTP changes
 - Safer for timing logic

Retry Calculation

- When the limit is exceeded, the server calculates how long the client must wait. This value becomes Retry-After.
 - Based on oldest request
 - Window size determines delay
 - Communicated to client

Implementing Token Bucket

- Token bucket refills tokens based on elapsed time. Each request consumes a token if available.
 - Track last refill time
 - Refill proportionally
 - Consume per request

Refill Rate

- The refill rate controls how quickly clients recover capacity. Faster refill allows higher sustained traffic.
 - Tokens per second
 - Controls steady-state rate
 - Independent of burst size

Burst Capacity

- Bucket capacity defines how many requests can be made instantly. This supports short bursts without overload.
 - Initial token count
 - Allows sudden spikes
 - Still rate-limited long-term

Flask Integration

- Rate limiting logic runs inside request handlers or hooks. It executes after authentication but before business logic.
 - Validate API key first
 - Apply rate limit check
 - Return early on failure

Returning 429 Responses

- When denying a request, the response must clearly explain why. HTTP 429 is the standard signal.
 - Status code 429
 - JSON error body
 - Retry-After header

Testing Behavior

- Testing rate limits requires sending many requests quickly. Curl loops are effective for this.
 - Rapid request loops
 - Observe status codes
 - Adjust sleep timing

Lab Completion

- At this point, you should understand how rate limiting works. The challenge applies it to your own API.
 - Per-key enforcement
 - Correct HTTP responses
 - Client-friendly behavior

Lab 2: Nginx as a Reverse Proxy

- In this lab, you place Nginx in front of a Flask web application. This mirrors how real production systems expose web apps safely.
 - Nginx faces the internet
 - Flask stays internal
 - Single entry point

What Is a Proxy?

- A proxy is an intermediary between a client and a server. It forwards requests and responses on behalf of another system.
 - Sits between client and server
 - Can modify requests
 - Can modify responses

Forward Proxy

- A forward proxy sits in front of clients. It controls how clients access the internet.
 - Client → Proxy → Internet
 - Controls outbound traffic
 - Common in corporate networks

Reverse Proxy

- A reverse proxy sits in front of servers instead of clients. Clients never talk to the backend directly.
 - Client → Proxy → Backend
 - Backends stay hidden
 - Clients see one endpoint

Why Reverse Proxies Exist

- Reverse proxies solve problems that applications should not handle directly. They centralize infrastructure concerns.
 - TLS termination
 - Security controls
 - Traffic management

Before and After Architecture

- This lab demonstrates how architecture changes when adding a reverse proxy. The backend becomes internal-only.
 - Before: Browser → Flask
 - After: Browser → Nginx → Flask
 - Flask no longer public

Nginx in This Lab

- Nginx will accept HTTPS traffic from the browser. It forwards requests to Flask over HTTP.
 - Listens on port 443
 - Proxies to localhost:5000
 - Handles TLS

Why Flask Should Not Face the Internet

- Flask's built-in server is for development only. Production traffic should be handled by proper servers.
 - Not hardened for attacks
 - Single-process model
 - No TLS management

Reverse Proxy Responsibilities

- The proxy becomes responsible for many cross-cutting concerns.
Applications can stay simple.
 - TLS certificates
 - Logging and headers
 - Rate limiting and auth

Lab Goal

- By the end of this lab, you will access Flask only through Nginx. Direct backend access should no longer be used.
 - HTML served through proxy
 - API served through proxy
 - Backend stays private

Flask Behind a Proxy

- When Flask runs behind Nginx, it no longer handles direct client connections. It becomes an internal service.
 - Listens on localhost only
 - No public exposure
 - Simpler application logic

Plain HTTP for Backend

- The backend Flask app runs over HTTP, not HTTPS. TLS is terminated at the reverse proxy.
 - HTTPS handled by Nginx
 - Flask runs without SSL
 - Clear separation of concerns

Jinja2 Rendering

- Flask renders HTML templates using Jinja2. Nginx simply forwards requests and responses.
 - Templates rendered by Flask
 - HTML returned to proxy
 - Proxy passes content unchanged

Serving HTML Through Nginx

- The browser only communicates with Nginx. HTML pages appear the same as before.
 - Same URLs
 - Same responses
 - Different architecture

Proxying API Requests

- JSON API endpoints are also proxied through Nginx. Frontend and API share the same origin.
 - /api/data via proxy
 - Backend still on :5000
 - Same-origin behavior

Why CORS Is Avoided

- Because Nginx serves both frontend and API from the same origin, browsers do not enforce CORS.
 - Same scheme
 - Same host
 - Same port

Headers Added by Nginx

- Nginx forwards important metadata to the backend using headers. Flask can read these headers.
 - X-Forwarded-For
 - X-Real-IP
 - X-Forwarded-Proto

Trusting Proxy Headers

- Backends must trust headers only from known proxies. Blind trust is dangerous.
 - Proxy controls headers
 - Do not trust direct clients
 - Secure internal network

Observing Client IP

- Flask can see the real client IP via forwarded headers. This enables logging and security controls.
 - `request.headers`
 - `X-Forwarded-For`
 - Proxy visibility

Lab Progress Check

- At this stage, traffic flows through Nginx to Flask. Direct backend access should be avoided.
 - Proxy path works
 - Backend internal only
 - Ready for production patterns

Nginx Configuration File

- Nginx behavior is defined by configuration files. A server block defines how requests are handled for a domain.
 - server { } block
 - listen and server_name
 - location routing

TLS Configuration

- Nginx handles HTTPS using certificates and private keys. This centralizes TLS management.
 - `ssl_certificate`
 - `ssl_certificate_key`
 - Single TLS termination point

proxy_pass Directive

- The proxy_pass directive forwards requests to the backend service. It defines where traffic is sent.
 - Points to backend URL
 - Supports HTTP backends
 - Transparent to client

Preserving Host Header

- Forwarding the Host header keeps backend routing consistent.
Some apps rely on it.
 - proxy_set_header Host
 - Supports virtual hosts
 - Avoids surprises

Forwarded IP Headers

- Nginx forwards client IP information using headers. Backends use these for logging and security.
 - X-Forwarded-For
 - X-Real-IP
 - X-Forwarded-Proto

Reloading Nginx Safely

- Configuration changes should be validated before reload. This prevents downtime.
 - nginx -t
 - systemctl reload nginx
 - Zero-downtime reload

Gunicorn and Flask

- Gunicorn is a WSGI server used to run Flask in production. It handles concurrency properly.
 - Multiple worker processes
 - Binds to localhost
 - Designed for production

Why Not Flask Dev Server

- The Flask development server is not production-ready. It lacks security and performance features.
 - Single process
 - No request hardening
 - Debug-only features

Production Request Flow

- In production, traffic flows through multiple layers. Each layer has a clear responsibility.
 - Browser → Nginx
 - Nginx → Gunicorn
 - Gunicorn → Flask app

Lab 2 Completion

- You now understand how reverse proxies fit into real deployments.
This mirrors cloud architectures.
 - Backend hidden
 - TLS centralized
 - Scalable pattern

Pop Quiz 1

- What is the primary role of a reverse proxy?
 - A. Hide clients from servers
 - B. Hide servers from clients
 - C. Replace backend apps
 - D. Encrypt databases



Pop Quiz 2

- Where does TLS typically terminate in this lab?
 - A. Flask app
 - B. Browser
 - C. Nginx
 - D. Gunicorn



Pop Quiz 3

- Why does this setup avoid CORS issues?
 - A. Flask disables CORS
 - B. Same-origin requests
 - C. Nginx rewrites headers
 - D. Browser ignores CORS



Pop Quiz 4

- Which header carries the original client IP?
 - A. Host
 - B. X-Forwarded-Proto
 - C. X-Forwarded-For
 - D. User-Agent



Pop Quiz 5

- Why should Flask not face the internet directly?
 - A. Flask is too slow
 - B. Flask lacks TLS support
 - C. Dev server is not production-ready
 - D. Flask cannot handle JSON



Answer 1

- A. Hide clients from servers
- **B. Hide servers from clients**
- C. Replace backend apps
- D. Encrypt databases



Answer 2

- A. Flask app
- B. Browser
- **C. Nginx**
- D. Gunicorn



Answer 3

- A. Flask disables CORS
- **B. Same-origin requests**
- C. Nginx rewrites headers
- D. Browser ignores CORS



Answer 4

- A. Host
- B. X-Forwarded-Proto
- **C. X-Forwarded-For**
- D. User-Agent



Answer 5

- A. Flask is too slow
- B. Flask lacks TLS support
- **C. Dev server is not production-ready**
- D. Flask cannot handle JSON



OAuth Login with GitHub

- OAuth lets users authorize your app without giving you their password. Your Flask app redirects to GitHub and receives a code to complete login.
 - Authorization, not password sharing
 - User can revoke access anytime
 - Common for “Login with ...” flows

Why OAuth Instead of Passwords

- With OAuth, your app never stores user passwords for the provider. Access is scoped and controlled by the provider.
 - No password handling
 - Least-privilege permissions
 - Revocable access

OAuth Roles

- OAuth defines four roles so everyone knows who does what. GitHub acts as both the authorization server and the API resource server.
 - Client: your Flask app
 - Resource Owner: the user
 - Authorization Server: GitHub login + token
 - Resource Server: GitHub API data

Authorization Code Flow

- The authorization code flow uses browser redirects to safely obtain an access token. The code is exchanged server-to-server.
 - Redirect user to GitHub
 - GitHub returns with ?code=...
 - Server exchanges code for token

Lab Endpoints

- This lab follows a simple set of routes for login, callback, and protected pages. Sessions store the logged-in user.
 - / (login page)
 - /login (redirect to GitHub)
 - /callback (exchange code for token)
 - /dashboard (protected)

Project Structure

- Keep templates separate from app logic. Your app file controls routes while templates render the pages.
 - app_oauth.py
 - requirements.txt
 - templates/layout.html
 - templates/login.html, dashboard.html

requirements.txt

- Only two packages are needed for this lab. Flask handles routing and sessions, and requests handles outbound HTTP calls.
 - Flask
 - requests
 - Install in a virtualenv

Register a GitHub OAuth App

- You must register an OAuth application so GitHub knows where to redirect users after approval. Callback URL must match exactly.
 - Homepage: <http://localhost:8000>
 - Callback: <http://localhost:8000/callback>
 - Copy Client ID and Client Secret

Environment Variables

- Store OAuth secrets in environment variables, not in code. This prevents accidental commits and makes configuration portable.
 - GITHUB_CLIENT_ID
 - GITHUB_CLIENT_SECRET
 - GITHUB_REDIRECT_URI

Templates Overview

- Templates render the UI while Flask handles logic. The login page links to /login and dashboard shows the username.
 - layout.html provides a base layout
 - login.html has “Login with GitHub” link
 - dashboard.html shows username + logout

OAuth Redirect Endpoint

- The login route builds a redirect URL that sends the user to GitHub's authorization server.
 - Includes client_id and redirect_uri
 - Specifies requested scopes
 - Generates a state value for CSRF protection

OAuth State Parameter

- The state parameter prevents CSRF by binding the auth request to the user session.
 - Generated randomly per login
 - Stored in session before redirect
 - Validated on callback

Authorization Code

- After approval, GitHub redirects back with a short-lived authorization code.
 - Code appears as query param
 - Code can be used only once
 - Expires quickly

Token Exchange

- The server exchanges the authorization code for an access token.
 - POST request to GitHub token endpoint
 - Client secret is used server-side
 - Browser never sees the token

Access Token Purpose

- The access token allows the app to call GitHub APIs on behalf of the user.
 - Sent in Authorization header
 - Limited by granted scope
 - Can be revoked by user

Fetching User Info

- The app uses the access token to fetch the authenticated user profile.
 - GET <https://api.github.com/user>
 - Returns username and id
 - Used to identify user locally

Creating a Session

- After login, user info is stored in the Flask session.
 - Session cookie sent to browser
 - Used for protected routes
 - No GitHub calls needed after login

Protecting Routes

- Protected routes check for authenticated session data.
 - Decorator pattern used
 - Redirects unauthenticated users
 - Matches real web apps

Logout Flow

- Logging out clears the local session but does not log out of GitHub.
 - Session is deleted server-side
 - User must re-authenticate
 - GitHub token remains valid

OAuth in Real Systems

- OAuth login is commonly used for SSO across many platforms.
 - GitHub, Google, Azure AD
 - Avoids password handling
 - Standardized and widely supported

Pop Quiz 1

- What problem does OAuth primarily solve?
 - A. Faster API responses
 - B. Password sharing between apps
 - C. Secure delegated access
 - D. Session encryption



Pop Quiz 2

- Which role verifies the user and issues tokens?
 - A. Client
 - B. Resource Owner
 - C. Authorization Server
 - D. Resource Server



Pop Quiz 3

- Why is the OAuth state parameter used?
 - A. Rate limiting
 - B. CSRF protection
 - C. Token encryption
 - D. Session expiration



Pop Quiz 4

- Which step exchanges the code for an access token?
 - A. /login redirect
 - B. Browser callback
 - C. Server-to-server POST
 - D. Dashboard load



Pop Quiz 5

- Where should OAuth access tokens be handled?
 - A. In browser JavaScript
 - B. In URL query strings
 - C. Server-side only
 - D. Local storage



Answer 1

- A. Faster API responses
- B. Password sharing between apps
- **C. Secure delegated access**
- D. Session encryption



Answer 2

- A. Client
- B. Resource Owner
- **C. Authorization Server**
- D. Resource Server



Answer 3

- A. Rate limiting
- **B. CSRF protection**
- C. Token encryption
- D. Session expiration



Answer 4

- A. /login redirect
- B. Browser callback
- **C. Server-to-server POST**
- D. Dashboard load



Answer 5

- A. In browser JavaScript
- B. In URL query strings
- **C. Server-side only**
- D. Local storage

